The Transformer Architecture

In 2017 researchers introduced the Transformer,

a model built entirely on self-attention.

It removed recurrence and allowed highly parallel training.

Scaling Laws

Around 2020 researchers observed that loss falls

as a power law of parameters, data and compute.

This observation motivated very large language models.